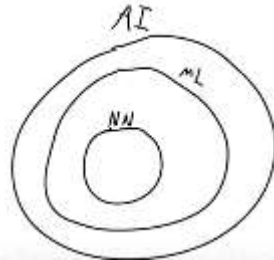
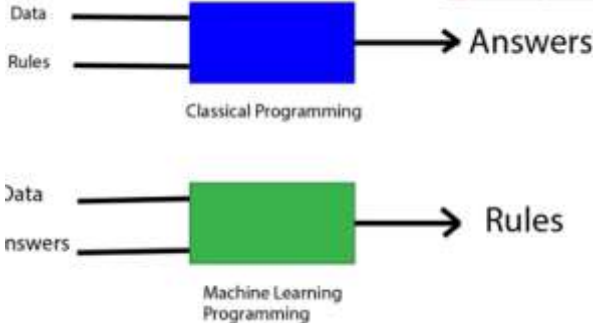
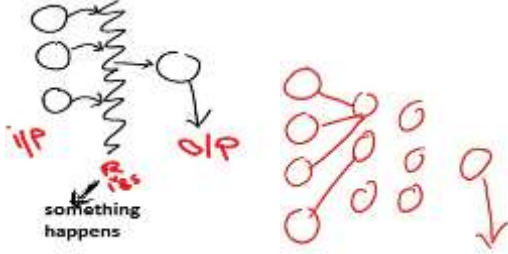



Machine Learning with Python

1.TensorFlow

- Machine Learning Fundamentals A - TensorFlow 2.0 Course - <https://www.youtube.com/watch?v=KwL1qTR5MT8>

	<p>AI - is the effort to automate intellectual tasks normally performed by humans</p> <ul style="list-style-type: none">• Early AI - used a predefined set of rules and coded into the computer.<ul style="list-style-type: none">○ No deep learning ML crazy algorithms happening.○ AI simply simulating some intellectual human behavior.• Now AI - has evolved into a much more complex field (ML, NN and other techniques)																
	<p>ML - is a part of AI. what is ML?</p> <ul style="list-style-type: none">○ AI used a predefined set of rules, means feed some data, go through the rules, analyze the data with the rules and produce the output○ Example of chess, in check, looks at the sets of rules and then it moves to somewhere else.• ML is actually figuring out the rules for us instead of hard coding the rules<ul style="list-style-type: none">○ Means take the input, output data and figure out the rules for us• ML requires a lot(ton) of input data to really train a good model.• ML models do not have 100% accuracy (trying to simulate like a human, can make mistakes), which means may not necessarily get the correct answer every single time.• Our goal create ML models is to raise accuracy as high as possible																
	<p>NN (Neural networks or deep learning) – It's a form of ML that uses a layered representation of data.</p> <ul style="list-style-type: none">• Input feed (bubbles) to this set of rules, something happens in here and get some output.• NN have more than 2(multiple) layers. An input layer (first layer of data), have some layers in between output layer, that are all connected together.																
 <table><tr><th></th><th>Feature</th><th>Label</th><th>Feature</th></tr><tr><td>1.</td><td>70</td><td>80</td><td>77</td></tr><tr><td>2.</td><td>60</td><td>90</td><td>84</td></tr><tr><td>3.</td><td>40</td><td>50</td><td>38</td></tr></table>		Feature	Label	Feature	1.	70	80	77	2.	60	90	84	3.	40	50	38	<p>Example, Students data set - info about students, midterm 1, 2 & final grade.</p> <ul style="list-style-type: none">• Student 1 - midterm 1 grade 70, midterm 2 grade 80 & final term grade 77• How can I use this information to predict any one of these three columns.<ul style="list-style-type: none">○ So if I were given a student's midterm 1 and final grade, how could I predict their midterm 2 grade.
	Feature	Label	Feature														
1.	70	80	77														
2.	60	90	84														
3.	40	50	38														

Features and labels

- **Features (Input information)** – Give midterm 1 and final grade to the model, to get some output is called features.
 - Training a model to look at midterm 1 and final grade, to make a new prediction
- **Labels (Output information – midterm 2)** trying to predict midterm 2.
 - Not have midterm 2 information & Not pass in the model, pass features (midterm 1 & final) to get the output of midterm 2

- **Intro to TensorFlow B - TensorFlow 2.0 Course** - <https://www.youtube.com/watch?v=r9hRyGGjOgQ>

- **what is a tensor?** It's a vector generalized to higher dimensions.
- **What is a vector?** Any linear algebra or basic vector calculus. It is kind of a data point. It doesn't necessarily have a certain coordinate.
- For example, **2 dimensional data point (x & y value)**,
- A vector can have any amount of dimensions (1, 2, 3, 4 (image data), 5 (video data))
- <https://www.tensorflow.org/guide/tensor> -
- **A tensor** is a generalization of vectors and matrices to potentially higher dimensions. Internally, TensorFlow represented tensors as n dimensional arrays of base datatypes.
- **Tensors**, are so important to **TensorFlow** - going to be working with manipulating and viewing

Each tensor represents a partially defined computation that will eventually produce a value.

TensorFlow is creating them & going to store partially defined computations in the graph.

Later, build the graph and have the session running, run different parts of the graph (execute different tensors) and get different results from our tensors.

Each tensor has a data type and a shape.

Data type is kind of information is stored in the tensor – like numbers, strings,

string = `tf.Variable("this is a string", tf.string)` – string tensor contains value and datatype

number = `tf.Variable(123, tf.int16)`

floating = `tf.Variable(123.456, tf.float64)`

Rank/Degree of tensors - the number of dimensions involved in the tensor.

`rank0_tensor = tf.Variable("first ", tf.string)` - **Scalar, contains a single value, and no "axes"**.

`rank1_tensor = tf.Variable(["first ", "OK"], tf.string)` - **Vector, a list of values & one axis.**

`rank2_tensor = tf.Variable([["first", "OK"], ["second", "yes"]], tf.string)` - **Matrix, has 2 axes**

`tf.rank(rank2_tensor)` => `<tf.Tensor: shape=(), dtype=string, numpy=2>` , **numpy=2 mean rank2**

Shapes of a tensor - how many items we have in each dimension.

Tensors have shapes. Some vocabulary:

Shape: The length (number of elements) of each of the axes of a tensor.

Rank: Number of tensor axes. A scalar has rank 0, a vector has rank 1, a matrix is rank 2.

Axis or Dimension: A particular dimension of a tensor.

Size: The total number of items in the tensor, the product of the shape vector's elements.

```
rank_0_tensor = tf.constant(4) => print(rank_0_tensor) =>  
tf.Tensor(4, shape=(), dtype=int32)
```

```
rank_1_tensor = tf.constant([2.0, 3.0, 4.0])
```


```
rank_2 = tf.constant([[1, 2], [3, 4], [5, 6]],  
dtype=tf.float16)
```

```
rank_3_tensor = tf.constant([  
    [[0, 1, 2, 3, 4], [5, 6, 7, 8, 9]],  
    [[10, 11, 12, 13, 14], [15, 16, 17, 18, 19]],  
    [[20, 21, 22, 23, 24], [25, 26, 27, 28, 29]],])
```


A scalar, shape: []

4




A vector, shape: [3]



A matrix, shape: [3, 2]



A 3-axis tensor, shape: [3, 2, 5]



```
rank_4 = tf.zeros([3, 2, 4, 5])
```

A tensor shape is like a vector. A 4-axis tensor

```
print("Type of every element:", rank_4.dtype)
print("Number of axes:", rank_4.ndim)
print("Shape of tensor:", rank_4.shape)
print("Elements along axis 0 of tensor:", rank_4.shape[0])
print("Elements along the last axis of tensor:",
rank_4.shape[-1])
print("Total number of elements (3*2*4*5): ",
tf.size(rank_4).numpy())
```

Type of every element: <dtype: 'float32'>

Number of axes: 4


Shape of tensor: (3, 2, 4, 5)

Elements along axis 0 of tensor: 3

Elements along the last axis of tensor: 5

Total number of elements (3*2*4*5): 120

A rank-4 tensor, shape: [3, 2, 4, 5]



Reshape/Changing/Manipulating shape

```
[ ] tensor1 = tf.ones([1,2,3]) # tf.ones() creates a shape [1,2,3] tensor full of ones
tensor2 = tf.reshape(tensor1, [2,3,1]) # reshape existing data to shape [2,3,1]
tensor3 = tf.reshape(tensor2, [3, -1]) # -1 tells the tensor to calculate the size of the dimension in that place
# this will reshape the tensor to [3,3]

# The number of elements in the reshaped tensor MUST match the number in the original
```

```
tf.Tensor(
[[[1. 1. 1.]
  [1. 1. 1.]]], shape=(1, 2, 3), dtype=float32)
tf.Tensor(
[[[1.]
  [1.]
  [1.]]

 [[1.]
  [1.]
  [1.]]], shape=(2, 3, 1), dtype=float32)
tf.Tensor(
[[[1. 1.]
  [1. 1.]
  [1. 1.]]], shape=(3, 2), dtype=float32)
```

```
%tensorflow_version 2.x
import tensorflow as tf
print(tf.version)

t = tf.zeros([5,5,5,5])
|
t = tf.reshape(t, [625])
```

```
rank_3_tensor = tf.constant([
[[0, 1, 2, 3, 4], [5, 6, 7, 8, 9]],
[[10, 11, 12, 13, 14], [15, 16, 17, 18, 19]],
[[20, 21, 22, 23, 24], [25, 26, 27, 28, 29]],])
```

```
print(rank_3_tensor)
```

```
tf.Tensor(
[[[ 0  1  2  3  4]
  [ 5  6  7  8  9]]

 [[10 11 12 13 14]
  [15 16 17 18 19]]

 [[20 21 22 23 24]
  [25 26 27 28 29]]], shape=(3, 2, 5),
dtype=int32)
```

A `-1` passed in the `shape` argument says "Whatever fits".

```
tf.Tensor(
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
 21 22 23 24 25 26 27 28 29], shape=(30,), dtype=int32)
```

Typically the only reasonable use of `tf.reshape` is to combine or split adjacent axes (or add/remove 1s).

For this 3x2x5 tensor, reshaping to (3x2)x5 or 3x(2x5) are both reasonable things to do, as the slices do not mix:

```
print(tf.reshape(rank_3_tensor, [3*2, 5]), "\n")
print(tf.reshape(rank_3_tensor, [3, -1]))
```


```
tf.Tensor(
[[[ 0  1  2  3  4] [ 5  6  7  8  9] [10 11 12 13 14]
  [15 16 17 18 19] [20 21 22 23 24]
  [25 26 27 28 29]], shape=(6, 5), dtype=int32)
```

```
tf.Tensor(
[[[ 0  1  2  3  4  5  6  7  8  9] [10 11 12 13 14 15 16 17 18 19]
```


```
print(tf.reshape(rank_3_tensor, [-1]))
```

```
[20 21 22 23 24 25 26 27 28 29]], shape=(3, 10),  
dtype=int32)
```

Some good reshapes.



Some bad reshapes.



Types of tensors – Variable, Constant, Placeholder and SparseTensor, a few other ones as well.

Except variable, all are immutable (value may not change during execution).

Evaluating tensors - create a session.

Sometimes need tensor object throughout our code, to do just use of default template

```
[ ] with tf.Session() as sess: # creates a session using the default graph  
    tensor.eval() # tensor will of course be the name of your tensor
```

- **Core Learning Algorithms A - TensorFlow 2.0 Course** - <https://www.youtube.com/watch?v=u5lZURgcWnU>

TensorFlow core learning algorithms, but not specific to TensorFlow, but they are used within there.

Linear regression

Classification

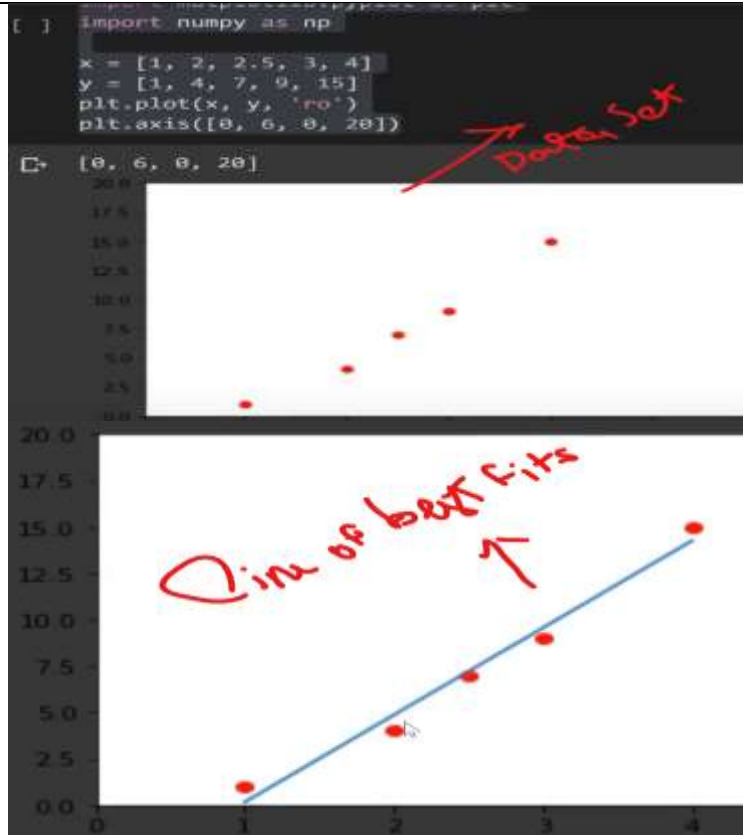
Clustering

Hidden Markov models


Now there is a ton time like 1000s of different machine learning algorithms.

Linear regression

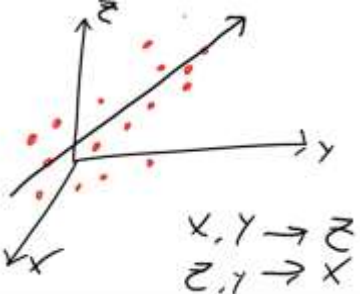
A linear correspondence between data points.



- Good example, plot a little graph & our data set
- use linear regression to come up with a model that can give us some good predictions for our data points.
- In this case, given some x value for a data point & predict the y value.
- See there's kind of some correspondence linearly for **these data points**.
- That means is draw something called a line of best fit through these data points
- Using this (**blue**) line, **can actually predict future values** in our data set.
- This is a very basic example for 2 dimensions with x and y. But oftentimes, have data points contains 8 or 9 kind of input values.
- **Line of best fit** refers to a line through a scatterplot of data points that best expresses the relationship between those points



- Use this line of best fit to predict a new data point.
- All red data points are trained our model with their information that gave to the model so that it could create this line of best fit
- **a line equation $\Rightarrow y = mx + b$**
- B stands for your y intercept (0.4)
- X and Y stands for the coordinates of this data point.
- M stands for the slope, which is probably the most important part.
- **Calculate the slope of a line** - draw a triangle, pick two data points, calculate 2 distance and divide the distance up by the distance across.
- Looks at all data points, line splits evenly. Means close to every data point as possible.

	<ul style="list-style-type: none"> • 2 dimension Equation $y = 1.5x + 0.5$ • X and Y don't have a value, that's because give the value (x or y) to come up with one (x or y) of the other ones (y or x). • If $x = 2$ then $y = 3.5$, data point as a prediction here on this line (black point). If $y = 2.7$, then find x.
	<ul style="list-style-type: none"> • 8 or 9 input variables, predict 1 output variable • 3 dimensions example, pass (x,y) -> predict z or (y,z) -> predict x
<p>Coding https://www.tensorflow.org/install</p> <p>install sklearn (even notebook) => “!pip install -q sklearn”</p> <p>install TensorFlow (notebook only) => “%tensorflow_version 2.x”</p> <p>from _future_import absolute_import, division, print_function, Unicode_literals</p> <p>import numpy as np import pandas as pd import matplotlib.pyplot as plt from IPython.display import clear_output from six.moves import urllib</p> <p>import tensorflow.compat.v2.feature_column as fc import tensorflow as tf</p>	<p>NumPy is a very optimized version of arrays in Python, for lots of multi dimensional array calculations.</p> <p>Pandas is a data analytics tool (loading/view/visualize data sets,...)</p> <p>Matplotlib is a for visual graphs and charts.</p> <p>The ipython display (notebook only) to clear the output.</p> <p>TensorFlow compact v2 feature column is for a feature column when we create a linear regression algorithm or model in TensorFlow</p> <p>Question Which type of analysis would be best suited for the following problem?:</p> <p>You have the average temperature in the month of March for the last 100 years. Using this data, you want to predict the average temperature in the month of March 5 years from now.</p>

Multiple regression
Correlation
Decision tree
[Answer]Linear regression

- **Core Learning Algorithms B - TensorFlow 2.0 Course** <https://www.youtube.com/watch?v=u85IOSsJsPI>

The Titanic data set - aimed to predict who's going to survive


```
import pandas as pd
dftrain = pd.read_csv("https://storage.googleapis.com/tf-datasets/titanic/train.csv")
```

survived	sex	age	n_siblings	parch	fare	class	deck	embark_t	alone
0	male	22	1	0	7.25	Third	unknown	Southamp	n
1	female	38	1	0	71.2833	First	C	Cherbourg	n
1	female	26	0	0	7.925	Third	unknown	Southamp	y
1	female	35	1	0	53.1	First	C	Southamp	n
0	male	28	0	0	8.4583	Third	unknown	Queenstov	y
0	male	2	3	1	21.075	Third	unknown	Southamp	n
1	female	27	0	2	11.1333	Third	unknown	Southamp	n
1	female	14	1	0	30.0708	Second	unknown	Cherbourg	n
1	female	4	1	1	16.7	Third	G	Southamp	n

```
dfeval = pd.read_csv("https://storage.googleapis.com/tf-datasets/titanic/eval.csv")
```


```
#print(dftrain.head()) #first 5 entries from data set
y_train = dftrain.pop('survived') #removes the column
y_eval = dfeval.pop('survived')
#print(dftrain.head())
#print(y_train)
#print(dftrain.loc[0], y_train.loc[0]) # locating row zero
#print(dftrain["age"])
#dftrain.describe() # describe the overall information
#dftrain.shape
#dftrain.age.hist(bins=20) # histogram of the age
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f196f2b6650>



```
#dftrain.sex.value_counts().plot(kind='barh')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f196f2a2050>



```
#dftrain['class'].value_counts().plot(kind='barh')
```

```
pd.concat([dftrain, y_train], axis=1).groupby('sex').survived.mean().plot(kind='barh').set_xlabel('% survive')
```

- **Core Learning Algorithms C - TensorFlow 2.0 Course** <https://www.youtube.com/watch?v=wz9J1slsi7I&t=6s>

2 different data - training (627,9) & testing (264, 9)

training data to create the model and testing data to evaluate and make sure that it's working properly.

doing machine learning models, typically have testing and training data.

categorical data is not numeric - **transform this data into numbers (integer, 0, 1,2), not in order**

numeric data – age

defined categorical & numeric columns – loop through them, create something called **feature columns**.

feature columns need to feed to our linear estimator or linear model to actually make predictions.

A CategoricalColumn with in-memory vocabulary -

https://www.tensorflow.org/api_docs/python/tf/feature_column/categorical_column_with_vocabulary_list?version=stable

Code

```
import pandas as pd
```

```
import tensorflow.compat.v2.feature_column as fc
```

```

import tensorflow as tf

dftrain = pd.read_csv("https://storage.googleapis.com/tf-datasets/titanic/train.csv")
dfeval = pd.read_csv("https://storage.googleapis.com/tf-datasets/titanic/eval.csv")
#y_train = dftrain.pop('survived') #removes the column
#y_eval = dfeval.pop('survived')

CATEGORICAL_COLUMNS = ['sex', 'n_siblings_spouses', 'parch', 'class', 'deck', 'embark_town', 'alone']
NUMERIC_COLUMNS = ['age', 'fare']

feature_columns = []
for feature_name in CATEGORICAL_COLUMNS:
    vocabulary = dftrain[feature_name].unique() # gets a list of all unique from given feature columns
    feature_columns.append(tf.feature_column.categorical_column_with_vocabulary_list(feature_name, vocabulary))

for feature_name in NUMERIC_COLUMNS:
    feature_columns.append(tf.feature_column.numeric_column(feature_name, dtype=tf.float32))

#print(feature_columns) #[VocabularyListCategoricalColumn(key='sex', vocabulary_list=('male', 'female'),
#dtype=tf.string, default_value=-1, num_oov_buckets=0),
#dftrain['sex'].unique() #array(['male', 'female'], dtype=object)

```

- **Core Learning Algorithms D - TensorFlow 2.0 Course** <https://www.youtube.com/watch?v=cEwvqVoBhI&t=1s>

Training Process - How do feed training data to the model?

In our case, only have 627 rows, it can fit that in PC RAM.

But if training a crazy ML model like 25 terabytes of data to pass it, can't load that into RAM

Batches - how this **process works**? Load, a small batch size of 32 entries at once to the model, that can increase our speed dramatically.

Don't load it entirely all at once, just load a specific set of kind of elements called **epochs**.

Epochs are essentially how many times the model is going to see the same data.

Overfitting – see(pass) the data too much to our model

Input function - is the way that we define how our data is going to be broke into epochs & into batches to feed model.

<https://www.tensorflow.org/tutorials/estimator/linear>

```
make_input_function(data_df, label_df, num_epochs=10, shuffle=True, batch_size=32): parameters
```

data_df - panda's data frame

label_df - labeled data frame (labels y train or eval)

num_epochs - number of epochs default 10

shuffle - shuffle data and mix it up before pass it to the model

batch_size - how many elements are we going to give to that to the model?

Code

```
from IPython.display import clear_output

def make_input_fn(data_df, label_df, num_epochs=10, shuffle=True, batch_size=32):
    def input_function(): # inner function, this will be returned
        #create a tf.data.Dataset object with the data and its label
        ds = tf.data.Dataset.from_tensor_slices((dict(data_df), label_df))
        if shuffle:
            ds = ds.shuffle(1000) # randomize order of data
        # split dataset into batches of 32 and repeat process for number of epochs.
        ds = ds.batch(batch_size).repeat(num_epochs)
        return ds #return a batch of the dataset
    return input_function # return a function object for use

train_input_fn = make_input_fn(dftrain, y_train) #call the function, returned a dataset object can feed to the model
eval_input_fn = make_input_fn(dfeval, y_eval, num_epochs=1, shuffle=False)

# Estimators are not recommended for new code. use keras API
# estimators are just basic implementations of algorithms in TensorFlow
linear_est = tf.estimator.LinearClassifier(feature_columns=feature_columns)

linear_est.train(train_input_fn) # train train_input_fn => train_input_fn()
result = linear_est.evaluate(eval_input_fn) # get model metrics/stats by testing on testing data

clear_output() #clear console output
print(result)
print(result['accuracy']) # the result variable is simply a dict of stats about our model
# 0.75757575 - This accuracy isn't very good, talk about how to improve this.

#re run
#Now notice, accuracy changed to 76 , the reason the data is getting shuffled and put in in a different order.
# It make different predictions and be trained differently.
#Change epochs, to 11, or 15, accuracy will change.
#goal is to get the most accurate model
```

#TensorFlow models are built to make predictions on a lot of things at once, they're not great at making predictions on
#like one piece of data (like one passenger to make a prediction for), they're much better at working in like large
#batches of data.
#Make a prediction for every single point that's in that evaluation data set.

#A dictionary that represents the predictions. For every single, what is it prediction.

#We passed 267 eval input data, it returned a list of all of these different dictionaries that represent each prediction.

```
pred_dicts = list(linear_est.predict(eval_input_fn))
print(pred_dicts)

#{'logits': array([-
2.0278394], dtype=float32), 'logistic': array([0.11631075], dtype=float32), 'probabilities': array([0.88
36892 , 0.11631081], dtype=float32),
#'class_ids': array([0]), 'classes': array([b'0'], dtype=object), 'all_class_ids': array([0, 1], dtype=i
nt32), 'all_classes': array([b'0', b'1'], dtype=object)}, ...

print(pred_dicts[0])
print(pred_dicts[0]['probabilities']) # [0.8836892 , 0.11631081] the percentage of survival 88% , won't
survive is 11%.
print(dfeval.loc[0]) # verify the passenger survey
print(y_eval.loc[0])
print(pred_dicts[0]['probabilities'][1]) # survival %
print(pred_dicts[0]['probabilities'][0]) # non survival %
```

Questions: What are epochs?

[Answer]The number of times the model will see the same data.

A type of graph.

The number of elements you feed to the model at once.

- **Core Learning Algorithms E - TensorFlow 2.0 Course** <https://www.youtube.com/watch?v=qFF7ZQNvK9E>

classification - is differentiating between data points and separating them into classes.

So rather than predicting a numeric value (did linear regression earlier), actually want to predict classes.
for example, iris flower data set, use different properties of flowers to predict what species of flower it is.
Iris flowers data - separates flowers into three different species (setosa, versicolor, virginica)

The information about each flower is (sepal/petal length & width, petal length & width) that information, Kara's - sub module of TensorFlow.

<https://www.tensorflow.org/tutorials/estimator/premade>

Code

```
import pandas as pd

import tensorflow.compat.v2.feature_column as fc
import tensorflow as tf

CSV_COLUMN_NAMES = ['SepalLength', 'SepalWidth', 'PetalLength', 'PetalWidth', 'Species']
SPECIES = ['Setosa', 'Versicolor', 'Virginica']

train_path = tf.keras.utils.get_file(
    "iris_training.csv", "https://storage.googleapis.com/download.tensorflow.org/data/iris_training.csv"
)
test_path = tf.keras.utils.get_file(
    "iris_test.csv", "https://storage.googleapis.com/download.tensorflow.org/data/iris_test.csv")

train = pd.read_csv(train_path, names=CSV_COLUMN_NAMES, header=0)
test = pd.read_csv(test_path, names=CSV_COLUMN_NAMES, header=0)
#here using Keras (a tensorflow module ) to grab datasets and read them into a pandas dataframe

train.head()

train_y = train.pop('Species')
test_y = test.pop('Species')

# The label column has now been removed from the features.
train.head()

train.shape # (120, 4)

#Input function

def input_fn(features, labels, training=True, batch_size=256):
    """An input function for training or evaluating"""
    # Convert the inputs to a Dataset.
    dataset = tf.data.Dataset.from_tensor_slices((dict(features), labels))

    # Shuffle and repeat if you are in training mode.
```



```

    if training:
        dataset = dataset.shuffle(1000).repeat()

    return dataset.batch(batch_size)

#Feature Columns
# Feature columns describe how to use the input.
my_feature_columns = []
for key in train.keys():
    my_feature_columns.append(tf.feature_column.numeric_column(key=key))

print(my_feature_columns)

```

Question : What is classification?

[Answer]The process of separating data points into different classes.

Predicting a numeric value or forecast based on independent and dependent variables.

None of the above.

- **Core Learning Algorithms F - TensorFlow 2.0 Course** <https://www.youtube.com/watch?v=5wHw8BTd2ZQ&t=9s>

Building the classification model – 100th of different classification pre-made models in TensorFlow.

Now can pick from, options are

DNNClassifier (deep neural network)

LinearClassifier – it works very similarly to linear regression, except it does classification.

Pick either one, but the DNN is the best choice, because may not be able to find a liner correspondence in our data.

It's not that difficult to change models, because most of the work comes from loading and pre processing our data.

Build a DNN with two hidden later with 30 nodes and 10 hidden nodes each.

hidden units is essentially us a building the architecture of the neural network.

An input layer, some middle layers (called hidden layers in a neural network), output layer

Decided 30 nodes in the 1st hidden layer 10 in the 2nd & the no of classes is 3 (3 classes for the flowers)

Code

```

# Build a DNN with 2 hidden layers with 30 and 10 hidden nodes each.
classifier = tf.estimator.DNNClassifier(
    feature_columns=my_feature_columns,
    # Two hidden layers of 30 and 10 nodes respectively.
    hidden_units=[30, 10],
    # The model must choose between 3 classes.
    n_classes=3)

```

Train the model

a lambda is an anonymous function that can be defined in one line

```
x = lambda: print("hi")
```

```
x1() # hi will print
```

steps = 5000 - similar to an epoch (go the dataset 10 times), but it go through the dataset until hit 5000 numbers
Run the code

it tells us the current step, the loss (the lower is the better), global steps per second.

Now at the end, final step, loss of 39, is pretty high, which means this is pretty bad.

Code

```
# Train the Model.
```

```
classifier.train(input_fn=lambda: input_fn(train, train_y, training=True), steps=5000)
```

Evaluation on the model.

Run, Much faster, we get a test accuracy of 80%.

Code

```
eval_result = classifier.evaluate(input_fn=lambda: input_fn(test, test_y, training=False))
```

```
print('\nTest set accuracy: {accuracy:0.3f}\n'.format(**eval_result))
```

Predictions on any given flower.

Type some numbers (petal length and width), then it will the predicted class of that flower is.

Not passing any y value, because we're making a prediction, so the model will the answer

type like 2.4, 2.6, 6.5, 6.3

notice we get three probabilities, one for each of the different classes.

class ID – it predicts is actually the flower, two means index array of 2

Code

```
import pandas as pd
```

```
#import tensorflow.compat.v2.feature_column as fc
```

```
import tensorflow as tf
```

```
CSV_COLUMN_NAMES = ['SepalLength', 'SepalWidth', 'PetalLength', 'PetalWidth', 'Species']
```

```
SPECIES = ['Setosa', 'Versicolor', 'Virginica']
```

```
train_path = tf.keras.utils.get_file(
```

```
    "iris_training.csv", "https://storage.googleapis.com/download.tensorflow.org/data/iris_training.csv"
```

```
)
```

```
test_path = tf.keras.utils.get_file(
```

```
    "iris_test.csv", "https://storage.googleapis.com/download.tensorflow.org/data/iris_test.csv")
```

```
train = pd.read_csv(train_path, names=CSV_COLUMN_NAMES, header=0)
```

```

test = pd.read_csv(test_path, names=CSV_COLUMN_NAMES, header=0)
#here using Keras (a tensorflow module ) to grab datasets and read them into a pandas dataframe

#train.head()

train_y = train.pop('Species')
test_y = test.pop('Species')

# The label column has now been removed from the features.
#train.head()

#train.shape # (120, 4)

#Input function

def input_fn(features, labels, training=True, batch_size=256):
    """An input function for training or evaluating"""
    # Convert the inputs to a Dataset.
    dataset = tf.data.Dataset.from_tensor_slices((dict(features), labels))

    # Shuffle and repeat if you are in training mode.
    if training:
        dataset = dataset.shuffle(1000).repeat()

    return dataset.batch(batch_size)

#Feature Columns
# Feature columns describe how to use the input.
my_feature_columns = []
for key in train.keys():
    my_feature_columns.append(tf.feature_column.numeric_column(key=key))

#print(my_feature_columns)

# Build a DNN with 2 hidden layers with 30 and 10 hidden nodes each.
classifier = tf.estimator.DNNClassifier(
    feature_columns=my_feature_columns,
    # Two hidden layers of 30 and 10 nodes respectively.
    hidden_units=[30, 10],
    # The model must choose between 3 classes.
    n_classes=3)

```

```

# Train the Model.
classifier.train(
    input_fn=lambda: input_fn(train, train_y, training=True),
    steps=5000)

eval_result = classifier.evaluate(
    input_fn=lambda: input_fn(test, test_y, training=False))

print('\nTest set accuracy: {accuracy:0.3f}\n'.format(**eval_result))

#Predict
#Not passing any y value, because model will y answer
def input_fn(features, batch_size=256):
    """An input function for prediction."""
    # Convert the inputs to a Dataset without labels.
    return tf.data.Dataset.from_tensor_slices(dict(features)).batch(batch_size)

features = ['SepalLength', 'SepalWidth', 'PetalLength', 'PetalWidth']
predict = {}

print("Please type numeric values as prompted:")
for feature in features:
    valid = True
    while valid:
        val = input(feature + ": ")
        if not val.isdigit(): valid = False

    predict[feature] = [float(val)]

predictions = classifier.predict(input_fn=lambda: input_fn(predict))

print(predictions)

for pred_dict in predictions:
    class_id = pred_dict['class_ids'][0]
    probability = pred_dict['probabilities'][class_id]

    print('Prediction is "{}" ({:.1f}%)'.format(SPECIES[class_id], 100 * probability))

```

```
WARNING:tensorflow:Using temporary folder as model direc
```

```
Test set accuracy: 0.867
```

```
Please type numeric values as prompted:
```

```
SepalLength: 2.3
```

```
SepalWidth: 2.6
```

```
PetalLength: 6.3
```

```
PetalWidth: 6.5
```

```
<generator object Estimator.predict at 0x7fb0f98d01d0>
```

```
Prediction is "Virginica" (92.2%)
```

Question: What kind of estimator/model does TensorFlow recommend using for classification?

LinearClassifier, [Answer]DNNClassifier, BoostedTreesClassifier

- **Core Learning Algorithms G - TensorFlow 2.0** Course <https://www.youtube.com/watch?v=8sqIaHc9Cz4&t=1s>

Clustering is the first unsupervised learning algorithm.

clustering only works for a very specific set of problems.

When use clustering, have a bunch of i/p information/features, don't have any labels or open information.

what clustering does is finds clusters of like data points and tells the location of those clusters.

Give a bunch of training data, pick how many clusters you want find.

Classifying handwritten digits using k means clustering. 10 different clusters for the digits 0 – 9.

The algorithm finds those clusters in the data set

Clustering

Now that we've covered regression and classification it's time to talk about clustering data!

Clustering is a Machine Learning technique that involves the grouping of data points. In theory, data points that are in the same group should have similar properties and/or features, while data points in different groups should have highly dissimilar properties and/or features.

(<https://towardsdatascience.com/the-5-clustering-algorithms-data-scientists-need-to-know-a36d136ef68>)


Unfortunately there are issues with the current version of TensorFlow and the implementation for KMeans. This means we cannot use KMeans without writing the algorithm from scratch. We aren't quite at that level yet so we'll just explain the basics of clustering for now.

Basic Algorithm for K-Means.

- Step 1: Randomly pick K points to place K centroids
- Step 2: Assign all of the data points to the centroids by distance. The closest centroid to a point is the one it is assigned to.
- Step 3: Average all of the points belonging to each centroid to find the middle of those clusters (center of mass). Place the corresponding centroids into that position.
- Step 4: Reassign every point once again to the closest centroid.
- Step 5: Repeat steps 3-4 until no point changes which centroid it belongs to.


Please refer to the video for an explanation of KMeans clustering.

Explain basic algorithm for K-Means




- 2 dimensions, make some data points (red), by putting them in their own unique little groups
- Now the algorithm starts for K means clustering. understand how this works, by randomly picking k centroids (filled in triangle) $k = 3$
- Now what happens next is each group/data point, is assigned to a cluster by distance.
- for every single data point find the distance using Euclidean or Manhattan distance,
- looking at this data point (circle) find the distance to all of these different centroids.


Assigned number (red above 1, 2, 3) to the closet to one data points




Added some more data points, now move these centroids into the middle of all of their data points called **center of mass**. Same thing other 2, remove and rearrange



keep doing until eventually reach a point where none of these points are changing the centroid, finally the draw is like



this now our cluster, if new points added, find the closet and assign to the closet cluster



Question : Which of the following steps is not part of the K-Means algorithm?

Randomly pick K points to place K centroids.

Assign each K point to the closest K centroid.

Move each K centroid into the middle of all of their data points.

[Answer] Shuffle the K points so they're redistributed randomly.

Reassign each K point to the closest K centroid.

- **Core Learning Algorithms H - TensorFlow 2.0 Course** <https://www.youtube.com/watch?v=lZg24y4wEPY&t=4s>

Hidden Markov models - deal with probability distributions.

- **Example weather model** - predict the weather on any day, given the probability of different events occurring.
- Using the average temperature on the days, create a hidden Markov model, will make a prediction for the weather in future days
- in this example, use some predefined probability distributions.

Hidden Markov Models

"The Hidden Markov Model is a finite set of states, each of which is associated with a (generally multidimensional) probability distribution []. Transitions among the states are governed by a set of probabilities called transition probabilities."

(<http://jedlik.phy.bme.hu/~gerjanos/HMM/node4.html>)

A hidden markov model works with probabilities to predict future events or states. In this section we will learn how to create a hidden markov model that can predict the weather.

https://www.tensorflow.org/probability/api_docs/python/tfp/distributions/HiddenMarkovModel

in a hidden Markov model, have a bunch of states.

Weather model, the states is hot & cold day (called **hidden**, because never access/look at these states)

In the model, called **observations** – each state have an observation

Example, If it is hot, 80% happy. If it is cold, 20% happy.

So at that state, we can observe the probability of something happening during that state is x or y

Data

Previous cases, use like 1000s of entries data points for our models to train for this.

Don't need any, only need is just constant values for probability(transition & observation distributions).

Data

Let's start by discussing the type of data we use when we work with a hidden markov model.

In the previous sections we worked with large datasets of 100's of different entries. For a markov model we are only interested in probability distributions that have to do with states.

We can find these probabilities from large datasets or may already have these values. We'll run through an example in a second that should clear some things up, but let's discuss the components of a markov model.

States: In each markov model we have a finite set of states. These states could be something like "warm" and "cold" or "high" and "low" or even "red", "green" and "blue". These states are "hidden" within the model, which means we do not directly observe them.

Observations: Each state has a particular outcome or observation associated with it based on a probability distribution. An example of this is the following: *On a hot day Tim has a 80% chance of being happy and a 20% chance of being sad.*

Transitions: Each state will have a probability defining the likelihood of transitioning to a different state. An example is the following: *a cold day has a 30% chance of being followed by a hot day and a 70% chance of being followed by another cold day.*

To create a hidden markov model we need.

- States
- Observation Distribution
- Transition Distribution

For our purpose we will assume we already have this information available as we attempt to predict the weather on a given day.


Draw Example - 2 state (s_1 , s_2)

probability of transitioning to the other state.

in a hot day, 20% chance of transitioning to a cold day, 80% chance of transitioning to another hot day

in a cold day, a 30% chance of transitioning to a hot day, 70% chance of transitioning to another cold day.

observation probability or distribution



Question : What makes a Hidden Markov model different than linear regression or classification?

[Answer] It uses probability distributions to predict future events or states.

It analyzes the relationship between independent and dependent variables to make predictions.

It separates data points into separate categories.